

# Data Scientist - Molecular Breeding -Interview

## Part 2: Genomic selection

By Zixiang Wen

### 1. Genomic selection for general combining ability (GCA) of Yield.

- 1.1 Try Lasso for yield
- 1.2 Try Ridge for yield
- 1.3 Try Random Forest for yield

### 2. Genomic selection for GCA of plant height.


- 2.1 Try Lasso for height
- 2.2 Try Ridge for height
- 2.3 Try Random Forest height

### 3. Genomic selection for GCA of Silk DAP.

- 3.1 Try Lasso for silk dap
- 3.2 Try Ridge for silk dap
- 3.3 Try Random Forest for silk dap

#### Major results from the above analysis:

1. Based on the above analysis, Random forest outperformed other methods in prediction accuracy at top 30% selection intensity across all three traits.



- 2. Prediction accuracy is higher for silk DAP which controlled by major gene (higher heritability), whereas lower for yield which controlled by multiple genes (lower heritability)
- 3. Mismatch between genotypic and phenotypic data hamper the proper sample size and prediction accuracy.

#### Works needed to be done in the future

- 1. Hyperparameter tuning for Random Forest.
- 2. Try Bayes method.
- 3. Enlarge population size for both genotyping and phenotyping
- 4. Separate Heterotic group for training model.

# Part 2 Genomic Selection Python Code and Output

November 26, 2019

## 0.1 1 Genomic selection for GCA of Yield

```
[1]: # main modules needed
import pandas as pd
import numpy as np
import seaborn as sns
from sklearn import linear_model
from sklearn.model_selection import train_test_split
from matplotlib import pyplot as plt
from scipy import stats
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import scale

[2]: num_g=pd.read_csv('Numeric data012_R.csv')

[3]: num_g.head()
```

	Unnamed: 0	2	3	4	8	9	10	11	12	13	...	303496	303499	303500	\
0		1	1	1	1	1	1	1	1	1	...	1	1	1	
1		2	0	2	0	0	1	1	2	0	0	...	1	1	1
2		3	1	1	1	1	1	1	1	1	...	1	1	1	
3		4	0	1	0	0	2	2	1	0	1	...	0	0	2
4		5	0	1	1	1	0	0	1	0	0	...	0	0	2

	303501	303502	303504	303505	303507	303510	303511
0	1	1	1	1	1	1	1
1	1	1	1	1	1	1	1
2	1	1	1	1	1	1	1
3	0	1	1	2	1	0	0
4	0	2	2	2	1	0	0

[5 rows x 136341 columns]

```
[4]: cname=pd.read_csv('c_name.csv')
[5]: num_g['Unnamed: 0']=cname['id']
[6]: num_g
```

```
[6]:      Unnamed: 0  2  3  4  8  9 10 11 12 13 ... 303496 303499 \
0      BLANK-1  1  1  1  1  1  1  1  1  1 ...    1    1
1      BLANK-2  0  2  0  0  1  1  2  0  0 ...    1    1
2      BLANK-3  1  1  1  1  1  1  1  1  1 ...    1    1
3      PHN11_0h43_0075  0  1  0  0  2  2  1  0  1 ...    0    0
4      W10004_0248  0  1  1  1  0  0  1  0  0 ...    0    0
...      ...  ...  ...  ...  ...  ...  ...  ...  ...  ...
1572     PI538011  0  1  1  0  0  0  0  1  1 ...    2    2
1573     Ames27138  0  2  0  0  2  0  0  0  0 ...    0    0
1574         Va35  0  2  0  0  2  1  1  0  0 ...    2    2
1575         Tx303  0  2  2  0  2  1  2  0  2 ...    0    2
1576     PI601773  0  2  0  0  0  0  0  1  1 ...    2    2
```

```
      303500  303501  303502  303504  303505  303507  303510  303511
0         1         1         1         1         1         1         1         1
1         1         1         1         1         1         1         1         1
2         1         1         1         1         1         1         1         1
3         2         0         1         1         2         1         0         0
4         2         0         2         2         2         1         0         0
...      ...      ...      ...      ...      ...      ...      ...      ...
1572        1         2         0         0         0         1         2         0
1573        0         0         2         2         0         0         0         0
1574        0         2         0         0         0         1         2         0
1575        0         2         2         2         0         2         0         0
1576        0         2         0         0         0         2         2         0
```

[1577 rows x 136341 columns]

```
[7]: num_g.rename(columns={"Unnamed: 0": "id"},inplace=True)
```

```
[8]: num_g.to_csv('clean_g_with-name.csv')
```

```
[9]: cn=pd.read_csv('clean_g_name.csv')
```

```
[10]: num_geno=pd.merge(num_g, cn,left_on=num_g.id, right_on=cn.id, how='inner')
```

```
[11]: num_geno=num_geno.drop(['key_0','id_y','simple_id'],axis=1)
```

```
[12]: num_geno.head()
```

```
[12]:      id_x  2  3  4  8  9 10 11 12 13 ... 303496 303499 \
0      AS6103  0  2  1  0  0  0  0  1  1 ...    0    0
1  PHN11_LH145_0029  0  2  0  0  0  0  0  0  0 ...    2    2
2      W10005_0107  0  2  0  0  0  0  0  0  0 ...    0    0
3      W10005_0032  0  2  0  0  1  1  0  0  0 ...    0    0
4      W10004_0082  1  1  1  1  1  0  0  1  1 ...    0    0

      303500  303501  303502  303504  303505  303507  303510  303511
0         0         0         2         2         0         0         0         0
1         0         2         0         0         0         2         2         0
```

2	2	0	2	2	0	1	0	0
3	2	0	2	2	0	0	0	0
4	0	0	2	1	2	1	0	0

[5 rows x 136341 columns]

```
[13]: num_genos.head()
```

```
[13]:
```

		id_x	2	3	4	8	9	10	11	12	13	...	303496	303499	\
0		AS6103	0	2	1	0	0	0	0	1	1	...	0	0	
1	PHN11_LH145_0029		0	2	0	0	0	0	0	0	0	...	2	2	
2	W10005_0107		0	2	0	0	0	0	0	0	0	...	0	0	
3	W10005_0032		0	2	0	0	1	1	0	0	0	...	0	0	
4	W10004_0082		1	1	1	1	1	0	0	1	1	...	0	0	

		303500	303501	303502	303504	303505	303507	303510	303511
0		0	0	2	2	0	0	0	0
1		0	2	0	0	0	2	2	0
2		2	0	2	2	0	1	0	0
3		2	0	2	2	0	0	0	0
4		0	0	2	1	2	1	0	0

[5 rows x 136341 columns]

```
[14]: p_y=pd.read_csv('y_blup_total.csv')
```

```
[15]: g_plus_p=pd.merge(num_genos,p_y, left_on=num_genos.id_x, right_on=p_y.id,
→how='inner')
```

```
[16]: g_plus_p.shape
```

```
[16]: (576, 136344)
```

```
[17]: g_plus_p
```

```
[17]:
```

		key_0	id_x	2	3	4	8	9	10	11	12	...	\
0	PHN11_LH145_0029	PHN11_LH145_0029	0	2	0	0	0	0	0	0	0	...	
1	W10005_0107	W10005_0107	0	2	0	0	0	0	0	0	0	...	
2	W10004_0082	W10004_0082	1	1	1	1	1	0	0	1	...		
3	PHN11_LH145_0028	PHN11_LH145_0028	0	2	0	0	0	0	0	0	0	...	
4	W10005_0029	W10005_0029	0	2	1	0	0	0	0	0	1	...	
...	...	...	...	...	...	...	...	...	...	...	...	...	
571	CM105	CM105	0	2	0	0	0	1	0	1	...		
572	Z013E0080	Z013E0080	0	0	0	1	2	1	0	0	...		
573	Z022E0009	Z022E0009	0	2	0	0	1	1	0	0	...		
574	Z022E0048	Z022E0048	0	2	0	0	1	1	0	0	...		
575	AH83	AH83	0	2	0	0	1	0	0	1	...		

		303500	303501	303502	303504	303505	303507	303510	303511	\
0		0	2	0	0	0	2	2	0	
1		2	0	2	2	0	1	0	0	

2	0	0	2	1	2	1	0	0
3	0	2	0	0	0	2	2	0
4	2	0	2	2	0	0	0	0
..	...	...	...	...	...	...	...	...
571	0	1	1	1	1	1	2	0
572	0	0	2	2	0	0	0	0
573	2	0	2	2	0	0	0	0
574	2	1	1	2	0	1	0	0
575	0	1	2	2	0	0	0	0

	id	yield
0	PHN11_LH145_0029	1.612589
1	W10005_0107	-3.249752
2	W10004_0082	1.901994
3	PHN11_LH145_0028	-10.440723
4	W10005_0029	6.706560
..	...	...
571	CM105	-6.000902
572	Z013E0080	8.373906
573	Z022E0009	-9.101649
574	Z022E0048	-2.145668
575	AH83	-51.816062

[576 rows x 136344 columns]

```
[18]: ## Since after inner join, part of data were discarded, it's necessary to check
      ↳ the outlier again.
```

```
Q1 = g_plus_p['yield'].quantile(0.25)
Q3 = g_plus_p['yield'].quantile(0.75)
IQR = Q3 - Q1
g_plus_p['yield_IQR'] = (g_plus_p['yield'] < (Q1 - 1.5 * IQR)) |
↳ (g_plus_p['yield'] > (Q3 + 1.5 * IQR))
```

```
[19]: from collections import Counter
```

```
[20]: Counter(g_plus_p['yield_IQR'])
```

```
[20]: Counter({False: 564, True: 12})
```

```
[21]: for i in range(g_plus_p.shape[0]):
      if g_plus_p['yield_IQR'][i]==True:
          g_plus_p['yield'][i]=np.nan
```

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel\_launcher.py:3:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

This is separate from the ipykernel package so we can avoid doing imports until

```
[22]: g_plus_p.shape
```

```
[22]: (576, 136345)
```

```
[23]: g_plus_p=g_plus_p.dropna()
```

```
[24]: X=g_plus_p.iloc[:,2:136342]  
Y=g_plus_p['yield']
```

```
[25]: # first trait analyzed  
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2)  
print(X_train.shape, y_train.shape)  
print(X_test.shape, y_test.shape)
```



```
(451, 136340) (451,)
```

```
(113, 136340) (113,)
```

```
[26]: print('      min   max   mean   sd')  
print('Train:', y_train.min(), y_train.max(), y_train.mean(), np.sqrt(y_train.  
    →var()))  
print('Test:', y_test.min(), y_test.max(), y_test.mean(), np.sqrt(y_test.var()))
```

```
      min   max   mean   sd  
Train: -25.04241474 16.96414301 -2.833114460982261 8.2296339886234  
Test: -24.65461518 19.30790655 -3.1429814772566367 8.181027222389439
```

```
[27]: plt.title('Train / test data')  
plt.hist(y_train, label='Train')  
plt.hist(y_test, label='Test')  
plt.legend(loc='best')  
plt.show()  
  
# marker PCA, use whole X with diff color for train and test  
X = np.concatenate((X_train, X_test))  
pca = PCA(n_components=2)  
p = pca.fit(X).fit_transform(X)  
Ntrain=X_train.shape[0]  
plt.title('PCA decomposition')  
plt.scatter(p[0:Ntrain,0], p[0:Ntrain,1], label='Train')  
plt.scatter(p[Ntrain:,0], p[Ntrain:,1], label='Test', color='orange')  
plt.legend(loc='best')  
plt.show()
```



```
[28]: ## simple model for GWAS  
pvals = []
```

```

for i in range(X_train.shape[1]):
    b, intercept, r_value, p_value, std_err = stats.linregress(np.
→asarray(X_train.iloc[:,i]), np.asarray(y_train))
    pvals.append(-np.log10(p_value))
pvals = np.array(pvals)
len(pvals)

```

[28]: 136340

```


[29]: # plot GWAS
plt.ylabel('-log10 P-value')
plt.xlabel('SNP')
plt.plot(pvals, marker='.')
plt.show()

# select by min_P_value
# min_P_value = 5 # P = 0.000001
# snp_list = np.nonzero(pvals>min_P_value)

## SNP selection based on p-value
N_best = 10000
snp_list = pvals.argsort()[-N_best:]

# finally slice X
X_train = X_train[X_train.columns[snp_list]]
X_test = X_test[X_test.columns[snp_list]]

```





### 0.1.1 1.1 Try Lasso for yield

```
[30]: ## try lasso
import sklearn.metrics as sm
# alpha is the regularization parameter
lasso = linear_model.Lasso(alpha=0.32)
lasso.fit(X_train, y_train)
y_hat = lasso.predict(X_test)


# mean squared error
mse = sm.mean_squared_error(y_test, y_hat)
print('\nMSE in prediction =',mse)

# correlation btw predicted and observed
corr = np.corrcoef(y_test,y_hat)[0,1]
print('\nCorr obs vs pred =',corr)

# plot observed vs. predicted targets
plt.title('Lasso: Observed vs Predicted Y')
plt.ylabel('Predicted')
plt.xlabel('Observed')
plt.scatter(y_test, y_hat, marker='.')
plt.show()
```

MSE in prediction = 51.51863639598614

Corr obs vs pred = 0.4802466586058425



```
[31]: ## Define the cutoff of 30% selection intensity

cutoff30=0.3*y_hat.shape[0]
print(cutoff30)

## Creat a function to calculate the accuracy under 30% selection intensity.
def ACC_30intensity (test,hat):
    yt=pd.DataFrame(test)
    yh=pd.DataFrame(hat)
    yh.index=yt.index
    combine=pd.concat([yt,yh],axis=1)
    a=pd.DataFrame(combine.iloc[:,0].argsort()[-34:])
    b=pd.DataFrame(combine.iloc[:,1].argsort()[-34:])
    ccc=pd.concat([a,b],axis=1)
    common=ccc.iloc[:,0].isin(ccc.iloc[:,1]).value_counts().to_frame().iloc[0,0]
    return print('Accuary under 30% selection intensity is',common/ccc.
    ↳shape[0])
```

33.9

```
[32]: ACC_30intensity (y_test,y_hat)
```

Accuary under 30% selection intensity is 0.5294117647058824

[33]: *## The above best alpha=0.42 was calculated from the following GridSearch CV*  
*→code:*

```
import warnings
warnings.filterwarnings('ignore')
from sklearn.linear_model import ElasticNet
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import GridSearchCV, train_test_split

# Create the hyperparameter grid
a = np.linspace(0, 1, 20)
param_grid = {'alpha': a}

# Instantiate the ElasticNet regressor: elastic_net
lasso = linear_model.Lasso()

# Setup the GridSearchCV
gm_cv=GridSearchCV(lasso,param_grid,cv=5)

# Fit it to the training data
gm_cv.fit(X_train,y_train)

# Predict on the test set and compute metrics
y_pred = gm_cv.predict(X_test)
r2 = gm_cv.score(X_test, y_test)
mse = mean_squared_error(y_test, y_pred)
print("Tuned lasso alpha: {}".format(gm_cv.best_params_))
print("Tuned R squared: {}".format(r2))
print("Tuned lasso MSE: {}".format(mse))
```

```
Tuned lasso alpha: {'alpha': 0.3157894736842105}
Tuned R squared: 0.21265587306975386
Tuned lasso MSE: 52.22997849603867
```

[34]: `from sklearn.model_selection import cross_val_score`  
`reg = lasso`  
*# Compute 10-fold cross-validation scores: cv\_scores*  
  
`cv_scores = cross_val_score(reg,X_train,y_train,cv=10)`  
  
*# Print the 10-fold cross-validation scores*  
`print(cv_scores)`  
  
`print("Average 10-Fold CV Score: {}".format(np.mean(cv_scores)))`

```
[ 0.09681262 -0.03737164  0.25932491  0.13584449  0.10822413  0.12529307
 0.15695921  0.13833454  0.02070628  0.05969437]
Average 10-Fold CV Score: 0.1063821981838089
```

### 0.1.2 1.2 Try Ridge for yield

```
[35]: ## Find best alpha for ridge regression
import warnings
warnings.filterwarnings('ignore')
# Import necessary modules
from sklearn.linear_model import RidgeCV
regr_cv = RidgeCV(alphas=[0.1, 1.0, 10.0])
model_cv = regr_cv.fit(X_train, y_train)
model_cv.alpha_
```

[35]: 10.0

```
[36]: ## Use the best alpha=10 to predict X_test

import sklearn.metrics as sm
# alpha is the regularization parameter
ridge = linear_model.Ridge(alpha=10)
ridge.fit(X_train, y_train)
y_hat = ridge.predict(X_test)


# mean squared error
mse = sm.mean_squared_error(y_test, y_hat)
print('\nMSE in prediction =',mse)

# correlation btw predicted and observed
corr = np.corrcoef(y_test,y_hat)[0,1]
print('\nCorr obs vs pred =',corr)

# plot observed vs. predicted targets
plt.title('Lasso: Observed vs Predicted Y')
plt.ylabel('Predicted')
plt.xlabel('Observed')
plt.scatter(y_test, y_hat, marker='.')
plt.show()
```

MSE in prediction = 56.66407953816876

Corr obs vs pred = 0.5179601793018527



```
[37]: ACC_30intensity (y_test,y_hat)
```

Accuary under 30% selection intensity is 0.5882352941176471

```
[38]: from sklearn.model_selection import cross_val_score
reg=ridge
# Compute 10-fold cross-validation scores: cv_scores
cv_scores = cross_val_score(reg,X_train,y_train,cv=10)
# Print the 10-fold cross-validation scores
print(cv_scores)

print("Average 10-Fold CV Score: {}".format(np.mean(cv_scores)))
```

```
[ 0.21884316 -0.13471916  0.37380859  0.26556722  0.24340188 -0.06092082
 0.27064002  0.17649123 -0.07587234  0.36677644]
Average 10-Fold CV Score: 0.16440162179797257
```

## 0.2 1.3 Try RandomForest Regressor for yield

```
[39]: from sklearn.ensemble import RandomForestRegressor
clf = RandomForestRegressor(120)
```

```
[40]: clf = RandomForestRegressor()
clf.fit(X_train, y_train)
```

```

y_hat = clf.predict(X_test)

# mean squared error
mse = sm.mean_squared_error(y_test, y_hat)
print('\nMSE in prediction =',mse)


# correlation btw predicted and observed
corr = np.corrcoef(y_test,y_hat)[0,1]
print('\nCorr obs vs pred =',corr)

# plot observed vs. predicted targets
plt.title('Random Forest: Observed vs Predicted Y')
plt.ylabel('Predicted')
plt.xlabel('Observed')
plt.scatter(y_test, y_hat, marker='o')
plt.show()

```

MSE in prediction = 46.34015635370706

Corr obs vs pred = 0.5518813878006141



[41]: ACC\_30intensity (y\_test,y\_hat)

Accuracy under 30% selection intensity is 0.5588235294117647

```
[42]: reg = clf

# Compute 10-fold cross-validation scores: cv_scores
cv_scores = cross_val_score(reg,X_train,y_train,cv=10,scoring='r2')

# Print the 10-fold cross-validation scores
print(cv_scores)
print("Average 10-Fold CV Score: {}".format(np.mean(cv_scores)))
```

```
[0.23649492 0.13182958 0.3171866  0.39873873 0.26170186 0.32231752
 0.35842361 0.30982434 0.20795259 0.30298462]
Average 5-Fold CV Score: 0.2847454372260335
```

### 0.3 2. Genomic selection for GCA of plant height

```
[43]: p_y_h=pd.read_csv('h_blup_total.csv')
```

```
[44]: g_plus_p=pd.merge(num_geno,p_y_h, left_on=num_geno.id_x, right_on=p_y_h.id,
→how='inner')
```

```
[45]: g_plus_p.shape
```

```
[45]: (575, 136344)
```

```
[46]: g_plus_p.head()
```

```
[46]:
```

	key_0	id_x	2	3	4	8	9	10	11	12	...	303500	\
0	PHN11_LH145_0029	PHN11_LH145_0029	0	2	0	0	0	0	0	0	...	0	
1	W10005_0107	W10005_0107	0	2	0	0	0	0	0	0	...	2	
2	W10004_0082	W10004_0082	1	1	1	1	1	0	0	1	...	0	
3	PHN11_LH145_0028	PHN11_LH145_0028	0	2	0	0	0	0	0	0	...	0	
4	W10005_0029	W10005_0029	0	2	1	0	0	0	0	1	...	2	

	303501	303502	303504	303505	303507	303510	303511	id	\
0	2	0	0	0	2	2	0	PHN11_LH145_0029	
1	0	2	2	0	1	0	0	W10005_0107	
2	0	2	1	2	1	0	0	W10004_0082	
3	2	0	0	0	2	2	0	PHN11_LH145_0028	
4	0	2	2	0	0	0	0	W10005_0029	

```

p_height
0 -2.085339
1 -13.954188
2 -2.117761
3 -15.784165
4  3.809174
```

```
[5 rows x 136344 columns]
```

```
[47]: X=g_plus_p.iloc[:,2:136342]
```

```
[48]: Y=g_plus_p['p_height']
```

```
[49]: # first trait analyzed  
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2)  
print(X_train.shape, y_train.shape)  
print(X_test.shape, y_test.shape)
```

```
(460, 136340) (460,)
```



```
(115, 136340) (115,)
```

```
[50]: print('      min   max  mean   sd')  
print('Train:', y_train.min(), y_train.max(), y_train.mean(), np.sqrt(y_train.  
→var()))  
print('Test:', y_test.min(), y_test.max(), y_test.mean(), np.sqrt(y_test.var()))
```

```
      min   max  mean   sd  
Train: -32.459111789999994 30.43450674 -1.8652208231717398 8.406251686869487  
Test: -32.459111789999994 21.11487856 -3.1070581186391304 9.092016742445761
```

```
[51]: plt.title('Train / test data')  
plt.hist(y_train, label='Train')  
plt.hist(y_test, label='Test')  
plt.legend(loc='best')  
plt.show()  
  
# marker PCA, use whole X with diff color for train and test  
X = np.concatenate((X_train, X_test))  
pca = PCA(n_components=2)  
p = pca.fit(X).fit_transform(X)  
Ntrain=X_train.shape[0]  
plt.title('PCA decomposition')  
plt.scatter(p[0:Ntrain,0], p[0:Ntrain,1], label='Train')  
plt.scatter(p[Ntrain:,0], p[Ntrain:,1], label='Test', color='orange')  
plt.legend(loc='best')  
plt.show()
```





```
[52]: pvals = []  
      for i in range(X_train.shape[1]):
```

```

    b, intercept, r_value, p_value, std_err = stats.linregress(np.
→asarray(X_train.iloc[:,i]), np.asarray(y_train))
    pvals.append(-np.log10(p_value))
pvals = np.array(pvals)
print(len(pvals))

```

136340

```


[53]: # plot GWAS
plt.ylabel('-log10 P-value')
plt.xlabel('SNP')
plt.plot(pvals, marker='o')
plt.show()

# select by min_P_value
#min_P_value = 5
#snp_list = np.nonzero(pvals>min_P_value)

N_best = 10000
snp_list = pvals.argsort()[-N_best:]

# finally slice X
X_train = X_train[X_train.columns[snp_list]]
X_test = X_test[X_test.columns[snp_list]]
snp_list

```



```
[53]: array([79424, 3161, 96743, ..., 6663, 19403, 25392], dtype=int64)
```

## 0.4 2.1 Try Lasso for height

```
[54]: ## Use the best a to predict X_test
import sklearn.metrics as sm
# alpha is the regularization parameter
lasso = linear_model.Lasso(alpha=0.26)
lasso.fit(X_train, y_train)
y_hat = lasso.predict(X_test)

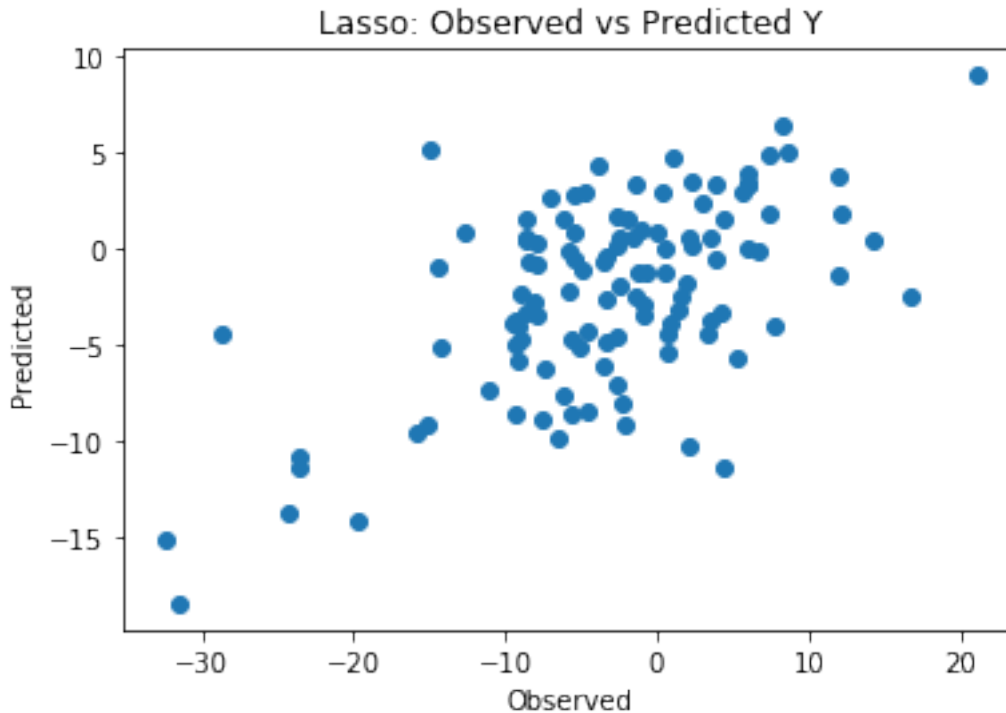
# mean squared error
mse = sm.mean_squared_error(y_test, y_hat)
print('\nMSE in prediction =',mse)

# correlation btw predicted and observed
corr = np.corrcoef(y_test,y_hat)[0,1]
print('\nCorr obs vs pred =',corr)

# plot observed vs. predicted targets
plt.title('Lasso: Observed vs Predicted Y')
plt.ylabel('Predicted')
plt.xlabel('Observed')
plt.scatter(y_test, y_hat, marker='o')
plt.show()
```

MSE in prediction = 53.67158851342249

Corr obs vs pred = 0.5974851858991395



```
[55]: ACC_30intensity (y_test,y_hat)
```

Accuary under 30% selection intensity is 0.5

```
[56]: ## Best alpha was calculated from the following code:
import warnings
warnings.filterwarnings('ignore')
# Import necessary modules
from sklearn.linear_model import ElasticNet
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import GridSearchCV, train_test_split

# Create the hyperparameter grid
a = np.linspace(0, 1, 20)
param_grid = {'alpha': a}

# Instantiate the ElasticNet regressor: elastic_net
lasso = linear_model.Lasso()

# Setup the GridSearchCV
gm_cv=GridSearchCV(lasso,param_grid,cv=5)

# Fit it to the training data
gm_cv.fit(X_train,y_train)
```

```

# Predict on the test set and compute metrics
y_pred = gm_cv.predict(X_test)
r2 = gm_cv.score(X_test, y_test)
mse = mean_squared_error(y_test, y_pred)
print("Tuned  alpha: {}".format(gm_cv.best_params_))
print("Tuned  R squared: {}".format(r2))
print("Tuned  MSE: {}".format(mse))

```

```

Tuned  alpha: {'alpha': 0.3684210526315789}
Tuned  R squared: 0.30354739551513654
Tuned  MSE: 57.07146638448894

```

```

[57]: reg = lasso
# Compute 10-fold cross-validation scores: cv_scores
cv_scores = cross_val_score(reg,X_train,y_train,cv=10,scoring='r2')

# Print the 10-fold cross-validation scores
print(cv_scores)

print("Average 10-Fold CV Score: {}".format(np.mean(cv_scores)))

```

```

[ 0.06008879  0.05643373  0.17290932  0.06738805  0.24109008 -0.02410407
  0.33645841  0.24498782  0.22654576  0.29071651]
Average 5-Fold CV Score: 0.1672514396025843

```

## 0.4.1 2.2 Try Ridge for height

```

[58]: ## Find best alpha for ridge regression
import warnings
warnings.filterwarnings('ignore')
# Import necessary modules
from sklearn.linear_model import RidgeCV
regr_cv = RidgeCV(alphas=[0.1, 1.0, 10.0])
model_cv = regr_cv.fit(X_train, y_train)
model_cv.alpha_

```

```

[58]: 10.0

```

```

[59]: ## try Ridge
import sklearn.metrics as sm
# alpha is the regularization parameter
ridge = linear_model.Ridge(alpha=10)
ridge.fit(X_train, y_train)
y_hat = ridge.predict(X_test)

# mean squared error
mse = sm.mean_squared_error(y_test, y_hat)

```

```

print('\nMSE in prediction =',mse)


# correlation btw predicted and observed
corr = np.corrcoef(y_test,y_hat)[0,1]
print('\nCorr obs vs pred =',corr)

# plot observed vs. predicted targets
plt.title('Lasso: Observed vs Predicted Y')
plt.ylabel('Predicted')
plt.xlabel('Observed')
plt.scatter(y_test, y_hat, marker='.')
plt.show()

```

MSE in prediction = 49.85027975535253

Corr obs vs pred = 0.6361211424442832



[60]: ACC\_30intensity (y\_test,y\_hat)

Accuary under 30% selection intensity is 0.5294117647058824

```
[61]: reg=ridge
# Compute 10-fold cross-validation scores: cv_scores

cv_scores = cross_val_score(reg,X_train,y_train,cv=10)

# Print the 10-fold cross-validation scores
print(cv_scores)

print("Average 10-Fold CV Score: {}".format(np.mean(cv_scores)))
```

```
[ 0.06698646  0.26677477  0.2024842   0.27302912  0.47938761 -0.15118739
  0.498393    0.15495287  0.37401065  0.47189177]
Average 5-Fold CV Score: 0.26367230602971437
```

#### 0.4.2 2.3 Try RandomForest Regressor for height

```
[62]: from sklearn.ensemble import RandomForestRegressor
clf = RandomForestRegressor(100)
```

```
[63]: clf.fit(X_train, y_train)
y_hat = clf.predict(X_test)


# mean squared error
mse = sm.mean_squared_error(y_test, y_hat)
print('\nMSE in prediction =',mse)

# correlation btw predicted and observed
corr = np.corrcoef(y_test,y_hat)[0,1]
print('\nCorr obs vs pred =',corr)

# plot observed vs. predicted targets
plt.title('Random Forest: Observed vs Predicted Y')
plt.ylabel('Predicted')
plt.xlabel('Observed')
plt.scatter(y_test, y_hat, marker='o')
plt.show()
```

MSE in prediction = 49.23582516560573

Corr obs vs pred = 0.6612758114785884



```
[64]: ACC_30intensity (y_test,y_hat)
```

Accuary under 30% selection intensity is 0.5

```
[65]: reg = clf
# Compute 10-fold cross-validation scores: cv_scores
cv_scores = cross_val_score(reg,X_train,y_train,cv=10,scoring='r2')

# Print the 10-fold cross-validation scores
print(cv_scores)

print("Average 10-Fold CV Score: {}".format(np.mean(cv_scores)))
```

```
[0.14692038 0.26536649 0.30613919 0.20222267 0.49720162 0.11871432
 0.4762445  0.39918695 0.45267231 0.49579089]
Average 10-Fold CV Score: 0.3360459328645949
```

### 0.5 3 Genomic selection for GCA of Silk DAP

```
[66]: p_y=pd.read_csv('s_blup_total.csv')
```

```
[67]: g_plus_p=pd.merge(num_geno,p_y, left_on=num_geno.id_x, right_on=p_y.id,
→how='inner')
```



```
[68]: g_plus_p.shape
```

```
[68]: (575, 136344)
```

```
[69]: g_plus_p.head()
```

```
[69]:
```

	key_0	id_x	2	3	4	8	9	10	11	12	...	303500	\
0	PHN11_LH145_0029	PHN11_LH145_0029	0	2	0	0	0	0	0	0	...	0	
1	W10005_0107	W10005_0107	0	2	0	0	0	0	0	0	...	2	
2	W10004_0082	W10004_0082	1	1	1	1	1	0	0	1	...	0	
3	PHN11_LH145_0028	PHN11_LH145_0028	0	2	0	0	0	0	0	0	...	0	
4	W10005_0029	W10005_0029	0	2	1	0	0	0	0	1	...	2	

	303501	303502	303504	303505	303507	303510	303511	id	\
0	2	0	0	0	2	2	0	PHN11_LH145_0029	
1	0	2	2	0	1	0	0	W10005_0107	
2	0	2	1	2	1	0	0	W10004_0082	
3	2	0	0	0	2	2	0	PHN11_LH145_0028	
4	0	2	2	0	0	0	0	W10005_0029	

	s_day
0	-0.443623
1	-1.722697
2	-0.732552
3	-0.425980
4	-1.471173

```
[5 rows x 136344 columns]
```

```
[70]: X=g_plus_p.iloc[:,2:136342]
```

```
[71]: Y=g_plus_p['s_day']
```

```
[72]: # first trait analyzed
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2)
print(X_train.shape, y_train.shape)
print(X_test.shape, y_test.shape)
```


```
(460, 136340) (460,)
(115, 136340) (115,)
```


```
[73]: print('      min    max    mean    sd')
print('Train:', y_train.min(), y_train.max(), y_train.mean(), np.sqrt(y_train.
→var()))
print('Test:', y_test.min(), y_test.max(), y_test.mean(), np.sqrt(y_test.var()))
```

	min	max	mean	sd
Train:	-10.69074918	8.077160751000001	0.0656371015771739	2.1012938122877243
Test:	-6.4046343325000015	4.1929703	-0.31234098288260875	2.221639448093306

```
[74]: plt.title('Train / test data')
plt.hist(y_train, label='Train')
plt.hist(y_test, label='Test')
plt.legend(loc='best')
plt.show()

# marker PCA, use whole X with diff color for train and test
X = np.concatenate((X_train, X_test))
pca = PCA(n_components=2)
p = pca.fit(X).fit_transform(X)
Ntrain=X_train.shape[0]
plt.title('PCA decomposition')
plt.scatter(p[0:Ntrain,0], p[0:Ntrain,1], label='Train')
plt.scatter(p[Ntrain:,0], p[Ntrain:,1], label='Test', color='orange')
plt.legend(loc='best')
plt.show()
```





```
[75]: ## Features selection based GWAS results
pvals = []
for i in range(X_train.shape[1]):
    b, intercept, r_value, p_value, std_err = stats.linregress(np.
→asarray(X_train.iloc[:,i]), np.asarray(y_train))
    pvals.append(-np.log10(p_value))
pvals = np.array(pvals)
print(len(pvals))
```


136340

```
[76]: # plot GWAS
plt.ylabel('-log10 P-value')
plt.xlabel('SNP')
plt.plot(pvals, marker='.')
plt.show()

# select SNP by min_P_value
# min_P_value = 12 #
# snp_list = np.nonzero(pvals>min_P_value)

N_best = 10000
snp_list = pvals.argsort()[-N_best:]
```

```
# finally slice X
X_train = X_train[X_train.columns[snp_list]]
X_test = X_test[X_test.columns[snp_list]]
leng=[item for item in snp_list]
```



### 0.5.1 3.1 Try Lasso Silk DAP

```
[77]: ## The best
## Cross-validation and greid search best a for lasso
# Create the hyperparameter grid
a = np.linspace(0, 1, 20)
param_grid = {'alpha': a}

# Instantiate the ElasticNet regressor: elastic_net
lasso = linear_model.Lasso()

# Setup the GridSearchCV
gm_cv=GridSearchCV(lasso,param_grid,cv=5)

# Fit it to the training data
gm_cv.fit(X_train,y_train)

# Predict on the test set and compute metrics
y_pred = gm_cv.predict(X_test)
```

```

r2 = gm_cv.score(X_test, y_test)
mse = mean_squared_error(y_test, y_pred)
print("Tuned lasso alpha: {}".format(gm_cv.best_params_))
print("Tuned lasso R squared: {}".format(r2))
print("Tuned MSE: {}".format(mse))

```

Tuned lasso alpha: {'alpha': 0.05263157894736842}  
 Tuned lasso R squared: 0.584461626437456  
 Tuned MSE: 2.03313072307695

```

[78]: ## Use the best alpha to predict the X-test
import sklearn.metrics as sm
# alpha is the regularization parameter
lasso = linear_model.Lasso(alpha=0.11)
lasso.fit(X_train, y_train)
y_hat = lasso.predict(X_test)

# mean squared error
mse = sm.mean_squared_error(y_test, y_hat)
print('\nMSE in prediction =',mse)


# correlation btw predicted and observed
corr = np.corrcoef(y_test,y_hat)[0,1]
print('\nCorr obs vs pred =',corr)

# plot observed vs. predicted targets
plt.title('Lasso: Observed vs Predicted Y')
plt.ylabel('Predicted')
plt.xlabel('Observed')
plt.scatter(y_test, y_hat, marker='.')
plt.show()

```

MSE in prediction = 2.388403219136664

Corr obs vs pred = 0.7302615281345879



```
[79]: ## Define the cutoff of 30%
cutoff30=0.3*y_hat.shape[0]
print(cutoff30)
## Creat a function to calculate the accuracy under 30% selection intensity.
def ACC_30intensity (test,hat):
    yt=pd.DataFrame(test)
    yh=pd.DataFrame(hat)
    yh.index=yt.index
    combine=pd.concat([yt,yh],axis=1)
    a=pd.DataFrame(combine.iloc[:,0].argsort()[-35:])
    b=pd.DataFrame(combine.iloc[:,1].argsort()[-35:])
    ccc=pd.concat([a,b],axis=1)
    common=ccc.iloc[:,0].isin(ccc.iloc[:,1]).value_counts().to_frame().iloc[0,0]
    return print('Accuary under 30% selection intensity is',common/ccc.
    ↳shape[0])
```

34.5

```
[80]: ACC_30intensity (y_test,y_hat)
```

Accuary under 30% selection intensity is 0.6857142857142857

```
[81]: training_accuracy = lasso.score(X_train, y_train)
test_accuracy = lasso.score(X_test, y_test)
print("Accuracy on training data: {:.2f}".format(training_accuracy))
print("Accuracy on test data:      {:.2f}".format(test_accuracy))
```

Accuracy on training data: 0.719402  
Accuracy on test data: 0.511850

```
[82]: # Compute 10-fold cross-validation scores: cv_scores

reg = lasso
cv_scores = cross_val_score(reg,X_train,y_train,cv=10,scoring='r2')

# Print the 10-fold cross-validation scores
print(cv_scores)

print("Average 10-Fold CV Score: {}".format(np.mean(cv_scores)))
```

[0.33042325 0.57912634 0.57720112 0.31091246 0.67373152 0.59992068  
0.32427942 0.69509696 0.35160778 0.51043897]  
Average 5-Fold CV Score: 0.49527385052655665

## 0.5.2 3.2 Try ridge for Silk DAP

```
[83]: ## Find best alpha for ridge regression
import warnings
warnings.filterwarnings('ignore')
# Import necessary modules
from sklearn.linear_model import RidgeCV
regr_cv = RidgeCV(alphas=[0.1, 1.0, 10.0])
model_cv = regr_cv.fit(X_train, y_train)
model_cv.alpha_
```

[83]: 10.0

```
[84]: ## try Ridge
import sklearn.metrics as sm
# alpha is the regularization parameter
ridge = linear_model.Ridge(alpha=10)
ridge.fit(X_train, y_train)
y_hat = ridge.predict(X_test)

# mean squared error
mse = sm.mean_squared_error(y_test, y_hat)
print('\nMSE in prediction =',mse)

# correlation btw predicted and observed
corr = np.corrcoef(y_test,y_hat)[0,1]
```

```


print('\nCorr obs vs pred =',corr)

# plot observed vs. predicted targets
plt.title('Lasso: Observed vs Predicted Y')
plt.ylabel('Predicted')
plt.xlabel('Observed')
plt.scatter(y_test, y_hat, marker='.')
plt.show()

```

MSE in prediction = 1.6396097653108253

Corr obs vs pred = 0.8193273719523999



[85]: ACC<sub>30</sub>intensity (y<sub>test</sub>,y<sub>hat</sub>)

Accuary under 30% selection intensity is 0.6571428571428571

```

[86]: reg=ridge
# Compute 10-fold cross-validation scores: cv_scores

cv_scores = cross_val_score(reg,X_train,y_train,cv=10)

```



```
# Print the 10-fold cross-validation scores
print(cv_scores)

print("Average 10-Fold CV Score: {}".format(np.mean(cv_scores)))
```

```
[ 0.42514713  0.70391462  0.62008691 -0.31770842  0.77879404  0.4841562
 0.11471732  0.61728037  0.14781385  0.31420127]
Average 5-Fold CV Score: 0.38884032971541754
```

### 0.5.3 3.3 Try Random Forest for Silk DAP

```
[87]: from sklearn.ensemble import RandomForestRegressor
      clf = RandomForestRegressor(100)

[88]: clf.fit(X_train, y_train)
      y_hat = clf.predict(X_test)


      # mean squared error
      mse = sm.mean_squared_error(y_test, y_hat)
      print('\nMSE in prediction =',mse)

      # correlation btw predicted and observed
      corr = np.corrcoef(y_test,y_hat)[0,1]
      print('\nCorr obs vs pred =',corr)

      # plot observed vs. predicted targets
      plt.title('Random Forest: Observed vs Predicted Y')
      plt.ylabel('Predicted')
      plt.xlabel('Observed')
      plt.scatter(y_test, y_hat, marker='.')
      plt.show()
```

MSE in prediction = 1.7729229680919807

Corr obs vs pred = 0.8207449081394613



```
[89]: ACC_30intensity (y_test,y_hat)
```

Accuracy under 30% selection intensity is 0.6857142857142857

```
[90]: training_accuracy = clf.score(X_train, y_train)
test_accuracy = clf.score(X_test, y_test)
print("Accuracy on training data: {:.2f}".format(training_accuracy))
print("Accuracy on test data:      {:.2f}".format(test_accuracy))
```

Accuracy on training data: 0.941146

Accuracy on test data: 0.637644

```
[91]: reg = clf
# Compute 10-fold cross-validation scores: cv_scores
cv_scores = cross_val_score(reg,X_train,y_train,cv=10,scoring='r2')

# Print the 10-fold cross-validation scores
print(cv_scores)

print("Average 10-Fold CV Score: {}".format(np.mean(cv_scores)))
```

```
[0.46121065 0.65952029 0.6810676  0.22180212 0.77829311 0.62204751
 0.27063762 0.72655111 0.49733918 0.56166547]
```

Average 5-Fold CV Score: 0.5480134657923644

```
[26]: acc=pd.read_csv('3t.csv')
```

```
[27]: sns.catplot(x='Method',y='Accuracy',hue='Accuracy_↵  
↪type',data=acc,col='Trait',kind='bar')
```

```
[27]: <seaborn.axisgrid.FacetGrid at 0xf0b00f0>
```

